

# Artificial Intelligence

Personal notes based on lecture material and assigned readings from Princeton's [COS 402: Artificial Intelligence](#), taught by Xiaoyan Li, Sebastian Seung, and Elad Hazan.

## Table of Contents

Tree and graph search .....	1
Adversarial search .....	3
Propositional logic.....	4
Bayesian networks .....	6
Temporal models .....	7
Markov decision processes .....	10

## Tree and graph search

- Algorithm properties
  - Completeness – an algorithm is complete if it terminates with a solution when one exists
  - Optimality – an algorithm is optimal if it finds the lowest path cost solution, where the path cost is (as of now) the sum of the individual step costs along the path
- Tree vs. graph search
  - Both maintain a *frontier set* of nodes under current consideration
  - Graph search maintains an *explored set* of nodes that have been visited
- Breadth-first search
  - Complete
  - Optimal if path cost is non-decreasing function of node depth
  - Parameters – branching factor  $b$ , depth  $d$
  - Time complexity is  $O(b^d)$
  - Space complexity is  $O(b^d)$ 
    - $O(b^{d-1})$  nodes in explored set and  $O(b^d)$  nodes in frontier
    - Limiting factor for breadth-first search
- Depth-first search
  - Tree-search version
    - Not complete in finite or infinite state spaces
    - Not optimal
    - Time complexity is  $O(b^m)$ , where  $m$  is the max depth of any node

- This value can be larger than the size of the state space
  - Space complexity is  $O(bm)$ 
    - Need only store single path from root to leaf node, plus unexpanded sibling nodes for each node on path
- Graph-search version
  - Complete in finite state space, but not in infinite
  - Not optimal
  - Time complexity bounded by size of state space
  - Space complexity – no advantage over BFS
- Depth-limited search
  - Not complete if depth limit  $l < d$
  - Not optimal if  $l > d$
  - Time complexity is  $O(b^l)$
  - Space complexity is  $O(bl)$
- Iterative deepening search
  - Combines benefits of breadth-first and depth-first search
  - Complete when branching factor is finite
  - Optimal when path cost is nondecreasing function of node depth
  - Space complexity is  $O(bd)$
  - Time complexity is  $O(b^d)$ , asymptotically the same as breadth-first search
- Bidirectional search
  - Uses breadth-first search in both directions (to and from goal state)
  - Time complexity is  $O(b^{d/2})$
  - Space complexity is  $O(b^{d/2})$
- Best-first search
  - Greedy best-first search
    - Uses  $f(n) = h(n)$
    - Not optimal
    - Completeness
      - Tree-search version is not complete, even in finite state spaces
      - Graph-search version is complete only in finite state spaces
    - Worst-case time and space complexity of  $O(b^m)$
  - A\* search
    - Uses  $f(n) = g(n) + h(n)$
    - Is both complete and optimal, given certain conditions (see below)
    - In most cases, number of states within goal contour search space is exponential in the length of the solution
    - Space complexity even worse, since it keeps all generated nodes in memory (like other graph-search algorithms)
- Search heuristic properties
  - Admissibility
    - A heuristic is admissible if it never overestimates cost to reach goal
  - Consistency

- A heuristic is consistent if it satisfies the triangle inequality – the heuristic value at node  $n$  is no greater than the heuristic value at node  $n'$  plus the step cost from  $n$  to  $n'$
- Facts
  - All consistent heuristics are admissible, but all admissible heuristics are not consistent
  - Tree-search version of A\* is optimal if  $h(n)$  is admissible
  - Graph-search version of A\* is optimal if  $h(n)$  is consistent
- Comparison of tree-search algorithms

Criterion	Breadth-First	Uniform-Cost	Depth-First <sup>5</sup>	Depth-Limited	Iterative Deepening	Bidirectional	A*
Complete?	Yes <sup>1</sup>	Yes <sup>1,2</sup>	No	No	Yes <sup>1</sup>	Yes <sup>1,4</sup>	Yes
Time	$O(b^d)$	$O(b^{1+\lceil C^*/\epsilon \rceil})$	$O(b^m)$	$O(b^l)$	$O(b^d)$	$O(b^{d/2})$	$O(b^{\epsilon d})$ <sup>8</sup>
Space	$O(b^d)$	$O(b^{1+\lceil C^*/\epsilon \rceil})$	$O(bm)$	$O(bl)$	$O(bd)$	$O(b^{d/2})$	Exponential <sup>9</sup>
Optimal?	Yes <sup>3</sup>	Yes	No	No	Yes <sup>3</sup>	Yes <sup>3,4</sup>	Yes <sup>6,7</sup>

<sup>1</sup> complete if branching factor  $b$  is finite

<sup>2</sup> complete if step costs  $\geq \epsilon$

<sup>3</sup> optimal if path cost is a non-decreasing function of depth  $d$

<sup>4</sup> if both directions use BFS

<sup>5</sup> graph-search DFS is complete in finite state spaces, and its space/time complexities are bounded by the size of the state space

<sup>6</sup> tree-search version is optimal if  $h(n)$  is admissible; graph-search version is optimal if  $h(n)$  is consistent

<sup>7</sup> A\* is *optimally efficient* – no other optimal algorithm is guaranteed to expand fewer nodes, except perhaps through tie-breaking among nodes with  $f(n) = C^*$  (on goal contour)

<sup>8</sup> bound holds in state space with single goal and constant step costs; extra cost is proportional to number of near-optimal goals

<sup>9</sup> see IDA\*, RBFS, (S)MA\*

## Adversarial search

- Minimax algorithm
  - Minimax value of a node is MAX's utility from being in that state, assuming both players play optimally
  - Minimax values are backed up through the tree from the leaves
  - Time complexity of  $O(b^m)$  and space complexity of  $O(bm)$  or  $O(m)$ 
    - Impractical for real games
- Alpha-beta pruning
  - Eliminates subtrees that could not possibly affect solution

## Propositional logic

- Terminology
  - If sentence  $\alpha$  is true in  $m$ , then  $m$  satisfies  $\alpha$  or  $m$  is a model of  $\alpha$
  - Entailment:  $\alpha$  entails  $\beta$  if and only if in every model in which  $\alpha$  is true,  $\beta$  is true, or,  $M(\alpha) \subseteq M(\beta)$
  - Logical equivalence: two sentences  $\alpha$  and  $\beta$  are logically equivalent if they are true in the same set of models
  - Validity: A sentence is valid if it is true in *all* models
  - Satisfiability: A sentence is satisfiable if its true in *some* model
- Basic results
  - Sentences  $\alpha$  and  $\beta$  are equivalent if and only if  $\alpha \models \beta$  and  $\beta \models \alpha$
  - $\alpha \models \beta$  if and only if the sentence  $\alpha \Rightarrow \beta$  is valid
  - $\alpha \models \beta$  if and only if  $\alpha \wedge \neg\beta$  is unsatisfiable
    - To prove that statement  $\alpha$  entails  $\beta$ , just show that  $\alpha \wedge \neg\beta$  is unsatisfiable
- Inference rules
  - Modus Ponens – if  $\alpha \Rightarrow \beta$  and  $\alpha$  holds, then  $\beta$  is true
  - And-Elimination – if  $\alpha \wedge \beta$  holds, then  $\alpha$  is true
- Resolution (algorithm)
  - Couples resolution (inference rule) with any complete search algorithm
  - Unit resolution rule
    - Take a clause and a literal and produces new clause
  - Full resolution rule
    - Takes 2 clauses (with complementary literals  $l_i$  and  $m_j$ ) and produces new literal containing all literals of 2 original clauses *except* the 2 complementary literals
    - Additionally, the resulting clause only contains one copy of each literal

$$\frac{l_1 \vee l_2 \vee \dots \vee l_k, \quad m_1 \vee m_2 \vee \dots \vee m_n}{l_1 \vee \dots \vee l_{i-1} \vee l_i \vee \dots \vee l_k \vee m_1 \vee \dots \vee m_{j-1} \vee m_{j+1} \vee \dots \vee m_n}$$

- Can be used to determine whether  $KB \models \alpha$ 
  - Strategy: show that  $KB \wedge \neg\alpha$  is unsatisfiable
  - Requires first converting  $KB, \alpha$  to conjunctive normal form
    - Clauses contain only ORed literals (or their negation)
    - Clauses are ANDed together
  - Resolution rule is applied to pairs of clauses from  $KB \wedge \neg\alpha$  with complementary literals
  - Process terminates when
    - No new clauses can be added –  $KB$  does not entail  $\alpha$
    - Two clauses resolve to the *empty* clause –  $KB$  entails  $\alpha$  (sentence is unsatisfiable)
      - Empty clause represents a contradiction (i.e.  $P \wedge \neg P$ )
  - Is complete – proved through ground resolution theorem

- Clause types
  - Definite clauses – disjunction of literals of which *exactly one* is positive
  - Horn clause – disjunction of literals of which *at most one* is positive
- Forward Chaining
  - FC-ENTAILS? ( $KB, q$ ) returns whether a knowledge base  $KB$  entails a query  $q$
  - Algorithm steps
    - Begins with set of known facts (positive literals) in knowledge base
    - If all premises of an implication are known, its conclusion is added to set of facts
    - Process continues until query  $q$  is added (returns true) or until no further inferences can be made (returns false)
  - Time complexity of  $O(n)$ , where  $n$  is the number of clauses
  - Is sound – every inference is an application of Modus Ponens
  - Is complete – every entailed atomic sentence will be derived (can be proved)
- Backward Chaining
  - Algorithm steps
    - Works backward from query  $q$
    - If  $q$  is true, no work needed (returns true)
    - Otherwise, finds implications whose conclusion is  $q$
    - Works backward until it reaches set of known facts
  - Time complexity of  $O(n)$  in theory, but often much less in practice
- DPLL
  - DPLL-SATISFIABLE? ( $s$ ) returns true or false
  - Recursive, depth-first search of possible models
  - Uses three heuristics
    - *Early termination* – if a literal is true, entire clause is true
    - *Pure symbol heuristic* – symbol that appears with sign in all clauses (value can be assigned, so as to make all clauses true)
    - *Unit clause heuristic* – a clause with a single literal (value must be assigned)
  - Various tricks can be used to improve performance
    - Intelligent backtracking (combined with conflict clause learning)
    - Random restarts
    - Clever indexing
- WalkSAT
  - WALKSAT( $clauses, p, max\_flips$ ) returns satisfying model or failure
  - On each iteration, algorithm picks unsatisfied clause and flips a symbol
    - Flips symbol that minimizes number of unsatisfied clauses or, with probability  $p$ , picks randomly
  - When algorithm returns failure, two possible causes
    - Sentence is unsatisfiable
    - Algorithm needs more time

## Bayesian networks

- Bayes Rule

$$P(b | a) = \frac{P(a | b)P(b)}{P(a)}$$

- Bayesian networks are a representation of the full joint probability

$$P(X_1, X_2, \dots, X_n) = \prod_i^n P(x_i | \text{Parents}(X_i))$$

- Conditional independence in Bayes Nets
  - Node is conditionally independent of non-descendants given its parents
  - Node is conditionally independent of *all* other nodes given its Markov blanket
    - Markov blanket consists of a node's parents, children, and spouses
- Exact inference
  - Enumeration
    - Uses  $P(X | e) = \alpha P(X, e) = \alpha \sum_y P(X, e, y)$  and full joint distribution to answer queries
    - ENUMERATION-ASK( $X, e, bn$ ) returns probability distribution  $P(X | e)$ 
      - Uses depth-first recursion (similar to DPLL for satisfiability)
      - Space complexity is  $O(n)$ , where  $n$  is the number of Boolean variables in the network
      - Time complexity is  $O(2^n)$
  - Variable elimination
    - Dynamic programming approach
    - Joint probability expressions are evaluated from right-to-left and intermediate results are stored
    - Uses matrices and pointwise product
    - Time and space complexity
      - Linear in size of the network for singly connected networks (polytrees)
      - Exponential for multiply connected networks
- Approximate inference
  - Direct sampling
    - Variables in probability expression are sampled in topological order
  - Rejection sampling
    - Generates samples from prior distribution specified by network, rejecting those that do not match evidence
    - Problem: rejects too many samples
      - Fraction of samples consistent with evidence  $e$  drops exponentially as number of evidence variables grows
  - Likelihood weighting
    - Fixes values for evidence variables  $E$

- Samples non-evidence variables, in any topological order
- Each event is weighted by its likelihood (the product of the conditional probabilities of each evidence variable, given its parents)
- Gibbs sampling (MCMC)
  - Evidence variables are fixed to their observed values
  - Nonevidence variables are randomly initialized
  - Generates a next state by randomly sampling for one of the nonevidence variables  $X_i$ 
    - Sampling is done conditioned on the current values of the variables in the Markov blanket of  $X_i$
  - Each state visited during process is a sample that contributes to the estimate for the query variable

## Temporal models

- Satisfy Markov assumption
  - Current state depends only on fixed, finite number of previous states
- Satisfy sensor (Markov) assumption
  - Current state is sufficient to determine current output (sensor) value
- Inference tasks
  - Filtering and prediction – computing belief state,  $P(X_t | e_{1:t})$ , or future state,  $P(X_{t+k} | e_{1:t})$ , given all evidence to date
    - Determining  $P(X_{t+1} | e_{1:t+1})$  from  $P(X_t | e_{1:t})$

$$P(X_{t+1} | e_{1:t+1}) = \alpha P(e_{t+1} | X_{t+1}) P(X_{t+1} | e_{1:t})$$

where

$$P(X_{t+1} | e_{1:t}) = \sum_{x_t} P(X_{t+1} | x_t) P(x_t | e_{1:t})$$

From this, define the FORWARD process as

$$f_{1:t+1} = \alpha \text{ FORWARD}(f_{1:t}, e_{t+1})$$

- Prediction is equivalent to filtering without evidence

$$P(X_{t+k+1} | e_{1:t}) = \sum_{x_{t+k}} P(X_{t+k+1} | x_{t+k}) P(x_{t+k} | e_{1:t})$$

- Smoothing – computing past state,  $P(X_k | e_{1:t})$ , given all evidence to date
  - Use forward message,  $f_{1:k}$ , and *backward* message,  $b_{k+1:t}$

$$\begin{aligned}
P(X_k | e_{1:t}) &= \alpha P(X_k | e_{1:k}) P(e_{k+1:t} | X_k) \\
&= \alpha f_{1:k} \times b_{k+1:t}
\end{aligned}$$

where

$$\begin{aligned}
b_{k+1:t} &= \sum_{x_{k+1}} P(e_{k+1} | x_{k+1}) P(e_{k+2:t} | x_{k+1}) P(x_{k+1} | X_k) \\
&= \text{BACKWARD}(b_{k+2:t}, e_{k+1})
\end{aligned}$$

- Time complexity of smoothing for a particular time step  $k$  with respect to evidence  $e_{1:t}$  is  $O(t)$
- To determine  $P(X_k | e_{1:t})$  for all  $1 \leq k \leq t$  in time  $O(t)$ , as opposed to time  $O(t^2)$ , and  $O(|f|t)$  space, where  $|f|$  is the size of the forward message, can use the *forward-backward algorithm*
  - FORWARD-BACKWARD(**ev**, *prior*) returns a vector of probability distributions  $P(X_k | e_{1:t})$  for  $1 \leq k \leq t$ 
    - Local variables
      - Vector of forward messages, **fv**
      - Representation of backward message, **b**
      - Vector of smooth estimates for steps 1...t, **sv**
    - Initialization
      - **fv**[0] =  $P(X_0)$
      - **b** =  $\vec{1}$
    - For  $i$  from 1 to  $t$ 
      - **fv**[ $i$ ]  $\leftarrow$  FORWARD(**fv**[ $i - 1$ ], **ev**[ $i$ ])
    - For  $i$  from  $t$  down to 1
      - **sv**[ $i$ ]  $\leftarrow$  NORMALIZE(**fv**[ $i$ ]  $\times$  **b**)
      - **b**  $\leftarrow$  BACKWARD(**b**, **ev**[ $i$ ])
    - Return **sv**
- Most likely explanation – given sequence of observations, find most likely sequence of states, i.e. compute  $\text{argmax}_{x_{1:t}} P(x_{1:t} | e_{1:t})$ 
  - Viterbi algorithm
    - Time complexity (like filtering) is  $O(t)$
    - Space complexity (unlike filtering) is  $O(t)$
    - Steps
      - Final maximization

$$\max_{x_{0:t}} P(x_{0:t}, e_{1:t}) = \max_{x_t} \left[ \max_{x_{0:t-1}} P(x_{0:t}, e_{1:t}) \right]$$



- Recursive step

$$\begin{aligned} & \max_{x_{0:t-1}} P(x_{0:t}, e_{1:t}) \\ &= \max_{x_{t-1}} \left[ P(x_t | x_{t-1}) P(e_t | x_t) \max_{x_{0:t-2}} P(x_{0:t-1}, e_{1:t-1}) \right] \end{aligned}$$

- Base case

$$\max_{x_{0:t-1}} P(x_{0:t}, e_{1:t}) = P(x_0)$$

- Hidden Markov Models (HMMs)
  - Temporal, probabilistic model in which state of process is described by a *single discrete* random variable
    - A model with 2 or more state variables can still be fit into the HMM framework by combining the variables into a single “megavariable”
  - Elegant matrix implementation of all basic algorithms possible
- Kalman filters
  - Filtering over continuous state variables
  - Required properties to do filtering
    - Probability distribution  $P(X_t | e_{1:t})$  is Gaussian
    - Transition model  $P(X_t | x_t)$  is linear Gaussian
    - Sensor model  $P(e_t | X_t)$  is linear Gaussian
    - Consequence: one-step predicted distribution will also be Gaussian

$$P(X_{t+1} | e_{1:t}) = \int P(X_{t+1} | x_t) P(x_t | e_{1:t}) dx_t$$

- If prediction  $P(X_{t+1} | e_{1:t})$  is Gaussian and sensor model  $P(e_{t+1} | X_{t+1})$  is linear Gaussian, then updated distribution is also Gaussian

$$P(X_{t+1} | e_{1:t+1}) = \alpha P(e_{t+1} | X_{t+1}) P(X_{t+1} | e_{1:t})$$

- Dynamic Bayesian Networks (DBNs)
  - Bayesian network that represents temporal probability model
  - DBNs and HMMs
    - Every HMM can be represented as a DBN with a single state variable and single evidence variable
    - Every discrete-variable DBN can be represented as an HMM
    - Advantage of DBNs
      - By decomposing a complex state into its constituent variables, DBNs can take advantage of sparseness in temporal model
      - A DBN with 20 Boolean state variables, each with 3 parents, will have  $20 \times 2^3 = 160$  probabilities in its transition model

- Corresponding HMM has  $2^{20}$  states or roughly a trillion probabilities in its transition matrix
- DBNs and Kalman filters
  - Every Kalman filter model can be represented as a DBN with continuous variables and linear Gaussian conditional distributions
  - However, not every DBN can be represented as a Kalman filter model
- Inference in DBNs
  - Possible approach: adapt variable elimination approach to filtering process, summing out state variables of previous time step to get distribution for new time step
    - Can achieve constant time and space per filtering update, but constant is exponential in number of state variables
  - Approximate inference solution: *particle filtering*
    - Algorithm steps
      - Initialize population of  $N$  samples by sampling from prior distribution  $P(X_0)$
      - For each time step
        - Propagate each sample forward by sampling next state value, given current, via  $P(X_{t+1} | x_t)$
        - Weight each sample by likelihood that it assigns to new evidence  $P(e_{t+1} | x_{t+1})$
        - Resample population to generate new population of  $N$  samples
          - New samples are selected from current population, with probability of selection proportional to the weight of the sample
    - Is consistent – gives correct probabilities as  $N$  tends to infinity
    - Efficient performance in practice

## Markov decision processes

- Key ideas
  - Discounted rewards
    - If discount factor  $\gamma$  is less than 1, utility of an infinite state sequence is still finite

$$U_h([s_0, s_1, s_2, \dots]) = R(s_0) + \gamma R(s_1) + \gamma^2 R(s_2) + \dots$$

- Expected utility
  - Expected utility obtained by executing a policy  $\pi$  starting in  $s$  is

$$U^\pi(s) = E \left[ \sum_{t=0}^{\infty} \gamma^t R(S_t) \right]$$

- Optimal policy

- Optimal policy is the policy with the highest expected utility

$$\pi_s^* = \operatorname{argmax}_{\pi} U^{\pi}(s)$$

- As a function of the state

$$\pi^*(s) = \operatorname{argmax}_{a \in A(s)} \sum_{s'} P(s' | s, a) U(s')$$

- Value iteration
  - Bellman equation for utilities

$$U(s) = R(s) + \gamma \max_{a \in A(s)} \sum_{s'} P(s' | s, a) U(s')$$

- Algorithm
  - Bellman update step

$$U_{i+1}(s) \leftarrow R(s) + \gamma \max_{a \in A(s)} \sum_{s'} P(s' | s, a) U_i(s')$$

- VALUE-ITERATION( $mdp, \epsilon$ ) returns a utility function
  - Local variables
    - $U, U'$  – vectors of utilities for states in  $S$ , initially zero
    - $\delta$  – max change in utility of any state in an iteration
  - Do while ( $\delta < \epsilon(1 - \gamma)/\gamma$ )
    - Initialization
      - $U \leftarrow U'$
      - $\delta \leftarrow 0$
    - For each state  $s$  in  $S$ 
      - $U'[s] \leftarrow R(s) + \gamma \max_{a \in A(s)} \sum_{s'} P(s' | s, a) U[s]$
      - If  $|U'[s] - U[s]| > \delta$ 
        - Set  $\delta \leftarrow |U'[s] - U[s]|$
  - Return  $U$

- Proof of convergence
  - Bellman update is a *contraction* by a factor of  $\gamma$  on the space of utility vectors with only 1 fixed point

- Policy iteration
  - Two main parts
    - Policy evaluation – given a policy  $\pi_i$ , compute the associated utility function  $U_i = U^{\pi_i}$ 
      - Uses simplified version of Bellman equation (policy is fixed)

$$U_i(s) = R(s) + \gamma \sum_{s'} P(s' | s, \pi_i(s)) U_i(s')$$

- Can do approximate policy evaluation by performing a fixed number  $k$  of simplified value iteration steps

$$U_{i+1}(s) \leftarrow R(s) + \gamma \sum_{s'} P(s'|s, \pi_i(s)) U_i(s')$$

- Policy improvement – calculate a new policy, using a one-step look-ahead based on  $U_i$
- Algorithm
  - POLICY-ITERATION( $mdp$ ) returns a policy
    - Local variables
      - $U$  – vectors of utilities for states in  $S$ , initially zero
      - $\pi$  – policy vector indexed by state, initially random
    - Repeat
      - Initialization
        - $U \leftarrow \text{POLICY-EVALUATION}(\pi, U, mdp)$
        - $\text{unchanged?} \leftarrow \text{true}$
      - For each state  $s$  in  $S$ 
        - If  $\max_a \sum_{s'} P(s'|s, a) U[s'] > \sum_{s'} P(s'|s, \pi[s]) U[s']$  then do
          - $\pi[s] \leftarrow \text{argmax}_{a \in A(s)} \sum_{s'} P(s'|s, a) U[s']$
          - $\text{unchanged?} \leftarrow \text{false}$
      - Break if  $\text{unchanged?} = \text{false}$
    - Return  $\pi$