

Information Security

Personal notes based on lecture material and assigned readings from Princeton's [COS 432: Information Security](#), taught by Ed Felten.

Table of Contents

	<u>Crypto Fundamentals</u>			<u>Firewalls and VPNs</u>
2	Message Integrity	15		Intranet
2	Pseudorandom Functions	15		NAT
2	Pseudorandom Generators	15		Firewalls
3	Encryption for Confidentiality	16		VPNs
3	Confidentiality and Integrity	16		DMZ
4	Pseudorandom Permutations			
4	Block Ciphers			<u>Web Security</u>
5	Encryption of Variable-Sized Messages	17		Web Browser
	• Cipher Modes	17		JavaScript
6	Asymmetric (Public) Key Cryptography	18		Cookies
6	RSA Algorithm	18		Cross-Site Request Forgery (CSRF)
7	Digital Certificates	19		Cross-Site Scripting (CSS/XSS)
8	Key Management	19		SQL Injection
8	Diffie-Hellman Key Exchange			
		20		Web Privacy
9	Password Security			• Tracking Techniques
10	Challenge-Response Protocols			• Anonymity?
10	Public Key Infrastructure	22		Electronic Voting
11	Access Control	23		Email Protocols
12	Zero-day Attack (Vulnerability)	23		Anonymous Communication
12	Secure Information Flow			• Tor
		24		SSH Protocol
	<u>Network Security</u>			
12	BGP and Shortest Path Routing			<u>Malware</u>
13	IP Packets and Spoofing	24		Taxonomy
	• DDoS Attack	24		Virus
13	Domain Name System (DNS)	25		Worms
	• "Cache Poisoning" Attack			
14	Cryptography in the Network Stack	26		Big Data and Privacy
		27		Economics and Security
14	POODLE and SSLv3 Vulnerability	28		Human Factors in Security

Important Concepts

Types of Security

- Confidentiality
- Integrity
- Availability

Message Integrity

- Use a MAC (message authentication code)
 - Alice sends Bob message m and $f(m)$
 - Bob receives (a, b) and accepts iff $f(a) = b$
- Properties of a secure MAC
 - Deterministic
 - Easily computable for Alice (to generate) and Bob (to verify)
 - Not easily computable by Mallory
- Solution: use PRF $f(k, m)$ (see below)
- Message order
 - Problem: Mallory could change message order or resend old messages
 - Solution: Append sequence number (nonce) to each message

Pseudorandom Functions (PRF)

- “As good as random” – indistinguishable from a random function that maps all inputs to 256-bit outputs generated by flipping 256 coins (truth table)
- Public family of functions f_0, f_1, f_2, \dots
 - Use f_k , where k is a secret key
- Theorem: If f is a PRF, then $f(k, m)$ with random(?) k is a secure MAC
 - Proof: If f is not a secure MAC, then f is not a secure PRF (contrapositive)
- Example: HMAC-SHA256

Pseudorandom Generators (PRG)

- Randomness often point of weakness in a security system
- PRGs use a small “seed” that is truly random
- Generate a long sequence of “good enough” (pseudorandom) values
 - Pseudorandom \cong unpredictable
 - PRG is *secure* if output indistinguishable from truly random values
- Important property: “hidden state”
- Desirable property: forward secrecy
 - If Mallory compromises hidden state of generator at time t , Mallory cannot backtrack to reconstruct past outputs of generator
- Examples
 - PRG that is secure, but lacks FS
 - init: (seed, 0)
 - adv: (seed, k) \rightarrow (seed, $k+1$)
 - out: $f(\text{seed}, k)$
 - PRG that is secure AND FS

- init: seed
- state S
- adv: $f(S, 0)$
- out: $f(S, 1)$
- Crux: state is overridden and f is not (feasibly) invertible

Encryption for Confidentiality

- Goal: ciphertext should not convey anything about plaintext
 - Semantic security (weaker than perfect secrecy)
 - Mallory chooses two plaintexts, we encrypt one of them
 - Mallory cannot do better than random guessing
- Alice encrypts with key k , Bob decrypts with key k
- First approach: one-time pad
 - Alice and Bob jointly determine long random string k (the pad)
 - Alice computes $E(k, x) = k \oplus x$
 - Bob computes $E(k, y) = k \oplus y = k \oplus (k \oplus x) = (k \oplus k) \oplus x = x$
 - Issues
 - Cannot reuse key
 - $(k \oplus a) \oplus (k \oplus b) = (a \oplus b)$
 - Easy to determine a and b with knowledge of English text distributions
 - Used in stream cipher attack (see below)
 - Need really long key (as long as sum of message lengths)
 - Strengths
 - Distribution of ciphertexts is random
 - Provably secure (Shannon 1949)
- Improvement: stream cipher
 - General idea: use PRG to “stretch” a small key into pseudorandom keystream
 - Start with truly random, fixed-size seed k
 - For each message, use a unique nonce (not necessarily secret)
 - Seed PRG with (k, nonce)
 - XOR message with output of PRG (like in one-time pad)
 - Critical: don’t reuse (k, nonce) pair!
 - Issues
 - Proof of security associated with one-time pad no longer holds

Confidentiality and Integrity

- Possible approaches
 - $E(x || M(x))$
 - Used by TLS/SSL (Transport Layer Security, Secure Sockets Layer)
 - Must decrypt ciphertext to check integrity (no integrity on ciphertext)
 - $E(x) || M(E(x))$
 - Used by IPSec (Internet Protocol Security) – winner!
 - Can determine integrity without decrypting ciphertext
 - $E(x) || M(x)$

- Used by SSH
 - Must decrypt ciphertext to check integrity
- Theorem: If E is semantically secure cipher and M is a secure MAC, then #2 is secure
 - Strategy of choice: encrypt plaintext, then append MAC of ciphertext
 - Bob first integrity checks, then decrypts
 - Important: Use separate keys for confidentiality and integrity
 - Important: Use separate pair of keys for Alice \rightarrow Bob and Bob \rightarrow Alice
 - This is authenticated encryption/decryption implemented in Assignment 1

Pseudorandom Permutations (PRP)

- Both encryption algorithm E and decryption algorithm D accept two inputs:
 - Block of size n bits
 - Key of size k bits
- Both E and D yield n -bit output block
- E_k is one of $(2^n)!$ permutations over the set of 2^n possible n -bit input blocks
- D is defined to be E^{-1}

	PR function	PR permutation	PR generators	Hash
Input	Any	Fixed-size	Fixed-size	Any
Output	Fixed-size	Fixed-size (equal)	Any	Fixed-size
Has key?	Yes	Yes	Yes (seed)	No
Invertible?	No	With key	No	Depends
Collisions	Yes, but can't find	No	No	Yes, but can't find

Block Ciphers

- Properties of a good block cipher
 - Efficiently computable (both E and $E^{-1} = D$)
 - Highly nonlinear ("confusion")
 - Hard for adversary to invert
 - Mix input bits together ("diffusion" / "avalanche effect")
 - Small changes in input create large/complicated changes in output
- Feistel network (type of block cipher)
 - Operates in d rounds, typically between 12 and 16
 - In each round i
 - Input is split into two halves, L_i and R_i
 - $L_{i+1} = R_i$
 - $R_{i+1} = L_i \oplus f(k_i, R_i)$
 - Final ciphertext: $R_d || L_d$ (no switch on last round)
 - Theorem: If f is a PRF, then 4-round feistel network is a PRP
- DES (Data Encryption Standard)
 - 64-bit blocks, 56-bit key
 - 16 (weak) Feistel rounds
 - History

- Designed in secrecy by IBM and NSA (1978)
 - U.S. government standard
 - Adopted by private sector
- Problems
 - Vulnerable to differential cryptanalysis (not publicly known)
 - Designed to be slow in software to discourage implementation
 - 56-bit key size: sufficient then, can be brute-forced now
- AES (Advanced Encryption Standard)
 - 128-, 192-, 256- bit versions (input, output, and key)
 - Ten rounds
 - Not feistel design
 - Symmetric-key algorithm
 - Same key used to encrypt and decrypt
 - Adopted by U.S. government and used worldwide (superseded DES)

Encryption of Variable-Sized Messages

- Padding
 - Plaintext not a multiple of blocksize
- Cipher modes
 - Multi-block messages
 - Schemes
 - ECB (Electronic Code Book)
 - Each block encrypted independently
 - Not semantically secure!
 - Does not hide data patterns
 - Subject to replay attacks
 - CBC (Cipher Block Chaining)
 - Each block of plaintext is XORed with previous ciphertext block before being encrypted
 - First block?
 - Generate random initialization vector (IV)
 - Treat as C_{-1} (prepend to final message)
 - XOR with first plaintext block before encryption
 - Decent solution
 - Identical messages -> different ciphertexts due to IV
 - Single bit errors propagates due to chaining
 - CTR (Counter mode)
 - Input to block cipher: messageID (nonce) || counter
 - Output of block cipher XORed with plaintext
 - Best solution!
 - Identical messages -> different ciphertexts due to nonce
 - Efficient to compute (parallelizable (enc/dec)ryption)
 - Note: If messageID is unique, can reuse key
 - Note: Not forward secret as a PRG

Asymmetric (Public) Key Cryptography

- Problems with symmetric key crypto
 - Integrity
 - Alice sending message to Bob, Charles, Diana (all share key k)
 - Bob can forge a message from Alice to Charles and Diana by computing MAC with shared key k
 - If n people are communicating, $\binom{n}{2}$ keys must be used
 - Confidentiality
 - Maybe only one party (e.g. Alice) should be able to decrypt message
 - Must exchange (secret) key in secure way
- Asymmetric scheme (idea first conceived by Diffie-Hellman-Cox in 1976)
 - Two different keys for encryption/decryption or signing (MAC)/verifying
 - One key is kept public and other is private

RSA Algorithm (Rivest-Shamir-Adleman 1978)

- Define $N = pq$, where p and q are large, randomly-chosen secret primes
- Pick e to be any value less than and relatively prime to $(p - 1)(q - 1)$
 - Can be small (3, 17, 65537 often chosen)
- Find d such that $ed \bmod (p - 1)(q - 1) = 1$
- Public key: (e, N)
 - e is the public key exponent
- Private key: (d, N) and (p, q)
 - d is the private key exponent
- Sending a message with RSA
 - Message M from Alice to Bob converted to integer $0 \leq m < n$
 - Use agreed upon, reversible padding scheme
 - Encryption: ciphertext $c \equiv m^e \pmod{N}$
 - Modular exponentiation is efficient
 - Decryption: plaintext $m \equiv c^d \pmod{N}$
- Relies on difficulty of integer factorization and “RSA problem”
 - Factorization
 - Factor modulus N to determine p, q
 - Compute $(p - 1)(q - 1)$
 - Determine d from $ed \bmod (p - 1)(q - 1) = 1$
 - RSA problem
 - Take e^{th} root of ciphertext modulo composite N
- Why is symmetric key crypto still used at all?
 - RSA is slow
 - Computationally weightier operations ($\sim 1000\times$ slower)
 - RSA keys are big
 - N is the product of two large primes ($\sim 4k$ bits)
- Applications
 - Confidentiality (“your eyes only” message)
 - Encrypt with public key of recipient
 - Recipient decrypts with private key

- Integrity (“digital signature”)
 - Sign by encrypting with private key
 - Verify by decrypting with public key
- Issues
 - Encrypting small messages with small e
 - If m^e is strictly less than N , ciphertext can be decrypted by taking e^{th} root of ciphertext
 - Chosen plaintext attack
 - Encrypt likely plaintexts under public key
 - Test if equal to ciphertext
 - Result: RSA not semantically secure
 - Why? RSA is a deterministic encryption algorithm
 - Chosen ciphertext attacks/malleability
 - Product of two ciphertexts is equal to encryption of product of respective plaintexts
 - Attacker asks private key holder to decrypt unsuspecting ciphertext $c' \equiv cr^e \pmod{N}$
 - r is chosen by the attacker
 - c' is the encryption of $mr \pmod{N}$
 - Attacker can multiply mr by r^{-1} to find m
 - Occasionally want a malleable cipher (not in RSA!)
 - Same plaintext \rightarrow same ciphertext (minor weakness)
 - RSA is deterministic
 - No built-in integrity check (minor weakness)
- Optimal Asymmetric Encryption Padding (OAEP)
 - Preprocessing step added before encryption to address all issues
 - Reverse OAEP used as a postprocessing step after decryption
 - Benefits
 - Adds element of randomness to deterministic RSA
 - Prevents partial decryption of ciphertexts/information leakage
- Encrypting larger messages
 - Cipher modes
 - CTR mode doesn't work because of randomization
 - CBC mode works, but is inefficient (overhead, very slow)
 - Hybrid encryption
 - Generate random symmetric encryption key k
 - Encrypt message with k
 - Encrypt k with RSA
 - Transmit both encrypted text and encrypted key

Digital Certificates

- Certifies the ownership of a public key by the named subject of the certificate
 - Address the impostor problem
- Chain of trust
 - Bob signs a message saying “Alice’s public key is ...”

- Works if Bob is known and believed to be trustworthy and competent
 - If we do not know, must ask Charlie to verify Bob's identity
- Certificate authority
 - Universally trustworthy third party lists verified public key holders
 - Everyone knows CA's public key
 - Customers of a CA are generally server administrators who need to present certificate to their clients

Key Management

- For symmetric ciphers, 128-bit keys are sufficient
- Need larger key for PFF/hash function
 - Finding a collision is more efficient than finding (exact) key
 - Birthday attack
 - If b is the bit-length of the hash, generate $2^{\frac{b}{2}}$ items at random
 - ~50% probability of finding collisions!
 - This attack takes only $O(2^{\frac{b}{2}})$ time and $O(2^{\frac{b}{2}})$ space (also possible in constant space)
 - PRF output size is typically 2x cipher output size (256 bits)
- Principles
 - Key management is usually the hard part
 - Keys must be strongly (pseudo)random
 - Each key should have a different purpose
 - Vulnerability of a key increases with
 - Usage
 - Places (copies) stored
 - Time
- Implications
 - Change keys periodically
 - Use "session keys"
 - Long-term keys used to negotiate session keys
 - Session key used temporarily
 - Erase keys when no longer needed
 - Keep keys out of long-term storage (if possible)
 - Keep keys in inaccessible places
 - Offline, locked in a safe
 - Protect against compromise of old keys (forward secrecy)

Diffie-Hellman Key Exchange (Diffie-Helman 1976)

- Relies on difficult of discrete-log problem (given $g^x \bmod p$, find x)
- Algorithm
 - Public: large prime p and primitive root g modulo p
 - p is often chosen to be $2q + 1$ where q is prime ("safe prime")
 - Alice selects random, secret a , $1 < a < p - 1$
 - Bob selects random, secret b , $1 < b < p - 1$

- Alice transmits $g^a \bmod p$ to Bob
- Bob transmits $g^b \bmod p$ to Alice
- Alice and Bob raise received values to their respective secret number
 - Arrive at shared secret $g^{ab} \bmod p$
- In practice: use $H(g^{ab} \bmod p)$ as shared secret key
- Insecure if adversary can modify messages (no authentication of communicating parties)
 - Man in the middle (MITM) attack
 - Instead of transmitting $g^a \bmod p$ to Bob, Mallory transmits $g^v \bmod p$
 - Instead of transmitting $g^b \bmod p$ to Alice, Mallory transmits $g^u \bmod p$
 - Alice ends up with $g^{au} \bmod p$ and Bob ends up with $g^{bv} \bmod p$
 - Mallory can forge messages between Alice and Bob by decrypting, modifying, and reencrypting
 - Solution: digital signature scheme
 - Server releases a public key to which it holds private key counterpart
 - Server signs hash of its copy of key with its private key
 - Client decrypts server signature with server's public key
 - If decryption matches hash of client's copy of key, then client concludes key is indeed shared
 - Client accepts messages from server if hashes match, otherwise closes channel
- Bad key values
 - If $g^a \bmod p$ or $g^b \bmod p$ is 1, shared key will be 1
 - If $g^a \bmod p$ or $g^b \bmod p$ is $p - 1$, shared key will be 1 or $p - 1$
 - Insecure values as adversary can guess key
 - Alice and Bob should reject if receive 1 or $p - 1$
 - Alice and Bob agree to reselect random number if mod power result is 1 or $p - 1$
 - Theorem: If p is a "safe prime" (see above), then 1 or $p - 1$ are the only insecure values
- Diffie-Hellman and forward secrecy
 - Alice and Bob have a shared key and want to negotiate new key
 - Alice and Bob conduct D-H key exchange protected by old key to get new key
 - Old key prevents adversary from modifying messages in D-H protocol
 - Alternative to using digital signature to *check* that keys match
 - If adversary *later learns* old key, cannot determine new key
 - D-H does not save or transmit information sufficient to determine key

Password Security

- Dictionary attack
 - Guessing attack using a precompiled list of likely options ("Password," "Computer," etc.)
 - Dictionary versus brute force attack
 - Brute force attack probes entire keyspace

- Brute force generally used against encryption, while dictionary attack is used against passwords (user generated)
- Storing the hash of a password in a database is better than storing the password
- Storing the *salted hash* of a password is better than storing just the hash
 - Two identical passwords will hash to the same value, so if an attacker cracks one password, the other is known to her as well
 - Solution: append a unique (not secret) value, called a *salt*, to a password before computing its hash
 - Store salt with the hash in the database
 - To verify a password, add salt to password, compute hash, and check against stored hash value
- Rainbow table attack
 - Precomputed table of hashes for possible plaintext passwords
 - Attack requires access to database of password hashes (“offline attack”)
 - Salts frustrate this attack, because for each possible password must compute hash corresponding to each salt in the database (adds dimension to table)

Challenge-Response Protocols

- Wish to authenticate user without revealing password p in protocol
- Procedure
 - User sends (user)name to server
 - Server asks user to encrypt random number r
 - User returns $PRF(p, r)$
 - Server verifies if user has correct password
- Used by HTTP (Web’s Hypertext Transfer Protocol)

Public Key Infrastructure

- Hardware/software/people/policies/procedures needed to create/manage/distribute/store/validate/revoke digital certificates
- Binds public keys with user identities through certificate authority (CA)
- Components of a PKI
 - Certificate authority (CA)
 - Root of trust in PKI, authenticating individuals, computers, network entities
 - CA issues own public key in self-signed CA certificate
 - Signs certificates with corresponding private key
 - Issues signed (encrypted) digital certificates
 - Alice requests certificate from CA
 - CA verifies Alice’s identity
 - CA computes hash of certificate contents and signs hash with own (CA’s) private key
 - CA appends signed hash to original certificate
 - CA makes Alice’s certificate publicly available
 - Verification process
 - Bob retrieves Alice’s certificate

- Bob decrypts signed hash with CA's public key
 - Bob compares hash of certificate with decryption
 - If match, knows Alice's public key is valid
- Registration authority
 - Role 1: subordinate CA
 - Certified by root CA to issues certificates for specific uses permitted by root
 - Role 2: verifies identity of users requesting info from CA
- Central database (server)
 - Holds certificate requests and record of issued/revoked certificates
- Certificate store (on local computer)
 - Saves issued certificates and record of pending/rejected requests
- Key Archival Server
 - Saves encrypted private keys in certificate database in case of loss
- Issues
 - Standards exist (X.509) but no government body enforcing standards
 - Provides chain of trust, but PKI is only as strong as weakest link
 - If one CA is compromise, security of entire PKI is at risk
 - 2011: Web browser vendors forced to blacklist all certificates issued by Dutch CA DigitNotar after 500 fake certificates discovered

Access Control

- SUBJECT wants to do VERB on OBJECT
 - Subject – active entity that requests access to an object or data within object
 - E.g. user, program, process, etc.
 - In this case, assume *running program*
 - Verb – action subject wishes to perform on object
 - In this case, assume *operation/API call*
 - Object – passive entity or resource that contains the information
 - E.g. computer, database, file, program, network connection, etc.
 - In this case, assume *system resource*
- Policy: set of allowed (S, V, O) triples
- Approaches
 - Access Control Matrix (ACM)
 - Table of Subjects v. Objects
 - Intersection contains allowed Verbs
 - Simple but inefficient implementation (matrix will be very sparse)
 - “Profiles”
 - For each user, store allowed permissions
 - Access Control List (ACL)
 - For each object, store (SUBJ, VERB) pairs
 - Alice: read, write; Bob: read; ...
 - Small and simple in practice
 - Most common approach (Parse!)

Zero-day Attack (Vulnerability)

- Exploitation of a previously unknown (to developers) vulnerability in a computer application or operating system
- Conducted in the time window between the discovery of the hole and the release of the security patch/update

Secure Information Flow

- Output of a program $P(v, s, r)$ should not leak any information about secret input s over all possible values of public/visible input v
- Cannot enforce by simply watching output (“dog that didn’t bark” problem)

Network Security

- Internet is a network of networks
 - Each network is an “autonomous system” (AS)
 - AS is a collection of routers (formally: IP routing prefixes)
 - Under the control of a single administrator, maintaining a clearly defined routing policy
 - AS’s connect together at exchange points
 - 47,000 AS numbers assigned by mid-2014
- Border Gateway Protocol (BGP)
 - Protocol designed to exchange routing and reachability information between autonomous systems (AS) on the Internet
 - Makes routing decision based on paths, network policies, and rule-sets configured by network administrator
 - Not actual routing protocol (that’s the Internet Protocol (IPv4, IPv6, ...))
- Shortest Path Routing
 - Simplified version of BGP
 - Attacks can lie about length of path to another router/host
 - Changes shortest path between routers A and B, diverting information through malicious router
 - Prefix hijacking attack
 - Also called IP hijacking or BGP hijacking
 - Involves announcing shorter route (either non-existent or tunneled) to redirect traffic
 - Pakistan’s attempt to censor Youtube (2008)
 - Accidentally leads to worldwide shutdown
 - China Telecom announces 37,000 prefixes not belong to them (2010)
 - Worldwide impact, but local traffic most affected
 - Malice *highly unlikely*
 - Can’t bypass application-level encryption
 - Can’t store all the traffic
 - Easily detectable
 - Defense
 - Cryptography can prevent an AS from lying about other nodes
 - Cannot prevent lying about their *own* links and costs

- Bottom line: relies on trust between small number of ASs
 - ASs can sever connections with a rogue node
 - Unlike application-layer security
- Layered network stack
 - Application
 - BGP – on top of TCP
 - Domain Name System (DNS) – on top of UDP
 - Transport (TCP or UDP)
 - Network (IP)
 - Physical and data link
- IP Packets and Spoofing
 - IP Packets contain source and destination IP addresses
 - Source can be spoofed, but destination can't
 - Return message will then be sent to spoofed address
 - Nodes cannot verify claimed source address
 - If A -> B -> C, C does not know if package originate from A or B
 - Node only knows local origin
 - Defenses
 - Ingress filtering
 - Discard an incoming packet if Source IP is inside network
 - Egress filtering
 - Discard an outgoing packet if Source IP is outside network
 - Distributed denial of service (DDoS) attack
 - Easy if lot of zombie nodes
 - Interesting: attack from single machine with bandwidth ~ as target
 - Smurf attack
 - Attacker sends broadcast ECHO request to network
 - Internet Control Message Protocol (ICMP) ping
 - Return address (source) is spoofed to be victim's address
 - All network hosts reply to victim
- Domain Name System (DNS)
 - Hierarchical, distributed naming system for devices connected to Internet
 - Translates domain names to IP addresses (DNS name resolution)
 - Process initiated on client side by DNS resolver
 - If a particular DNS server cannot translate, will ask another server
 - Recursive process
 - Caches recent translations
 - “Cache poisoning” attack (DNS spoofing)
 - Basic DNS does not use crypto
 - Attacker supplies incorrect translation of a domain name
 - Incorrect translation is cached (poisoning the cache)
 - Subsequent requests for translation of that domain name return address of server controlled by attacker
 - Solution: DNSSEC

- Provides: authentication of DNS data, authenticated denial of existence
 - Does not provide: availability or confidentiality (no encryption)
 - Answers from DNSSEC protected zones are digitally signed
 - DNS root servers used as root of trust
 - DNS hierarchy used as chain of trust
 - Parent domain (DNS zone) verifies DNSKey record in subdomains
 - Procedure
 - Domain owners generate their own keys
 - Upload them with DNS control panel to domain name-registrar
 - Keys pushed via secDNS to zone operator
 - Zone operator (i.e. Verisign for .com) signs and publishes keys in DNS
- Cryptography in the Network Stack
 - Can be incorporated into different layers (application level, transport level, network/IP level, etc)
 - SSL/TSL
 - Application layer (or between transport and application layers)
 - Authenticates: hostnames (server identity)
 - Server usually well-known entity
 - Encrypts: sessions over TCP layer
 - Allows secure communication (confidentiality and integrity) between server and client
 - IPsec
 - Network layer security
 - Goal: integrity/confidentiality at level of IP packets
 - Authenticates: IP addresses
 - Encrypts: IP packets
 - Problems
 - IP is stateless, but keeping state required for encryption
 - Communication consists of independent (request, response) pairs
 - Does not require server to retain session information over multiple requests
 - Many security problems are application-specific

POODLE and SSLv3 Vulnerability

- TSL/SSL are encryption protocols used to protect communication between websites and computers
 - Represented by small padlock icon in browser
 - Protects information from being intercepted, spied upon, or modified by attackers between user and service provider

- Prevents someone sharing Wi-Fi in Starbucks from spying on your bank transactions
 - TLS has now replaced SSL, except in cases of backward compatibility
- TLS clients will downgrade protocol used to lower version of TLS and then SSL if dealing with legacy servers (“downgrade dance”)
 - First handshake attempt: offers highest protocol version supported
 - If handshake fails, retry will earlier protocol versions
 - If attacker interferes with client-server negotiations, can downgrade to SSL 3
 - Attacker is MITM between client and server
- Encryption in SSL 3.0
 - Uses either RC4 stream cipher or block cipher in CBC mode
 - RC4 has biases
 - If same secret is sent over many connections and encrypted with many RC4 streams, information about secret will leak
 - CBC encryption
 - Block cipher padding is not deterministic, nor covered by MAC
 - Integrity of padding cannot be fully verified when decrypting
 - Attacker can decrypt “secure” HTTP cookies
- Solution
 - Disable SSL 3.0 in browser
 - Can prevent communication with legacy systems
 - Use TLS_FALLBACK_SCSV

Firewalls and VPNs

- Intranet
 - Private network internal to company
 - Private IP space (typically)
 - Internal view different from external view (<http://benefits/>)
 - Principle: don’t connect (most) machines directly to Internet
- Network Address Translation (NAT)
 - Machines assigned IPs from reserved spaces
 - Examples: 192.168.*.* and 10.*.*.*
 - Network shares single “real” IP address
 - NAT keeps translation table of inside IP address to outside IP equivalents
 - A router acts as an agent between the Internet and local network
 - No publicly visible IPS for local machines
 - Can’t accept incoming connections (directly)
- Firewalls
 - Perimeter defense for a network
 - Separate outside from inside
 - Monitor boundary
 - Block questionable incoming traffic
 - Centralize security policy for easy administration
 - Types
 - Network/IP layer: packet filtering

- Stateless packet filtering
 - Block all incoming connections
 - Stateful packet filtering
 - To block *some* incoming connections
 - Only allow incoming packet if in response to previous outgoing packet
 - Remember TCP sequence number and acknowledgement number
 - Allows more sophisticated policies
 - TCP layer: circuit-level gateway
 - Goal: allow servers to run on inside of firewall
 - Allows more sophisticated filtering than Network/IP layer
 - Components
 - SOCKS: TCP-level proxy protocol
 - Client library: internal machine
 - Client program: supports proxying
 - Application level: proxy server
 - Even more sophisticated filtering
 - Need separate proxies for each service
 - Complications
 - Firewall blocks incoming DNS replies
 - Use DNS proxy (application-level gateway)
 - Need to serve web and other content
 - Put servers outside firewall (DMZ)
 - Firewall blocks incoming email
 - Drop server outside firewall
 - Siphon mail in after filtering
- Virtual Private Networks (VPNs)
 - Extend the perimeter
 - Goal: make branch offices behave as if on same private network
 - VPN server is on the firewall
 - Dual interface with Internal IP and External IP
 - Example process
 - User working from home can use VPN client to entire private network
 - Authenticates to VPN server using username/password
 - Obtains shared session key
 - VPN server assigns intranet IP to client
 - Adds mapping of intranet IP – external IP to NAT table
 - “Tunnel” established
- DMZ
 - Firewall configuration used to secure local area networks (LANs)
 - Most computers run behind firewall connected to public network
 - One or more computers run outside firewall, in DMZ
 - Intercept traffic and broker request for rest of LAN
 - Another firewall separates these computers from rest of Internet

Web Security

- Browser (based on OS/hardware) interacts with website (based in network)
 - Browser sends requests
 - Website replies
- Two sides of web security
 - Web browser
 - Can be attacked by any website it visits
 - Attacks can lead to malware installation (keyloggers, botnets), document theft, loss of private data
 - Web application
 - Runs at website
 - Written in PHP, ASP, JSP, Ruby
 - Many potential bugs: CSRF, XSS, SQL injection
 - Attacks lead to stolen credit cards, defaced sites, etc.
- Web attacker
 - Entices user to visit malicious website (i.e. attacker.com)
 - Can easily obtain SSL/TLS certificate for his site (\$0)
 - Uses phishing email, enticing content, appears in search results, is placed by ad network, etc.
 - Network attacker
 - Passive: wireless eavesdropper
 - Active: evil router, DNS poisoning (see network security notes)
 - Malware attacker
 - Attacker control's user machine
 - How? Convinces user to install malicious content
 - Masquerades as antivirus program, codec for new video format, etc.
 - Exploits application bugs (e.g. buffer overflow)
- JavaScript
 - Language executed by browser
 - Scripts embedded in Web pages
 - Can run before HTML is loaded, before page is viewed, while it is being viewed, or when user is leaving the page (any time)
 - Used to implement "active" web pages
 - AJAX (asynchronous JavaScript and XML) allows Web apps to send/retrieve data from server in background
 - Note: JSON often used instead of XML, need not be asynchronous
 - Origin of many security issues
 - Allows attacker to execute code on user's machine (browser)
 - Security model
 - Script runs in "sandbox"
 - No direct file access (on user's computer)
 - Can cause browser to load remote pages/resources like scripts or images, which may be cached locally by browser

- Can technically only store cookies
- Same-origin policy
 - Can only read properties of documents and windows from same *server, protocol, and port*
 - Does not apply to library imports
 - Scripts loaded in enclosing frame from external site
 - `<script type = "text/javascript" src="http://www...">`
`</script>`
 - Script runs as if loaded from site that provided page
- Cookies
 - Small piece of data sent from website and stored in user's browser
 - Every time user loads website, browser sends cookie back to server
 - Purpose: allow websites to remember state (items in a shopping cart) and/or record user's browsing activity (clicking buttons, logging in, etc)
 - Can store form content such as: passwords, credit card number, address
- Three attacks
 - Cross-Site Request Forgery (CSRF)
 - Malicious script (from user's visit to malicious site) makes forged request to "good" site with user's cookie for good site
 - Changes Netflix accounts settings, steals Gmail contacts
 - At risk: web applications that perform actions on input from authenticated users without requiring authorization of specific action
 - User authenticated by cookie
 - Browser tricked into sending HTTP request to target site
 - Can force user to perform state changing requests:
 - Transferring funds out of bank account
 - Changing email address/password
 - *Cannot* directly see website's response to browser's forged request
 - Attacker must find URL that has side effects or online form
 - Unless...attacker uses cross-site scripting
 - Defenses
 - Secret validation token (synchronizer token pattern)
 - Secret and unique token embedded by web app into all HTML forms and verified on server side
 - Must ensure unpredictability and uniqueness (i.e. using hash chain of random seed)
 - Example: `<input type="hidden" name="..." value="Kby...">`
 - Can be difficult on web apps that heavily use AJAX
 - Referer validation
 - Check HTTP Referer header to ensure request is coming from authorized page
 - May cause issues with browsers that omit Referer header for privacy reasons (too strict a policy)

- Must calibrate leniency of policy
 - E.g. Referer: `http://www.facebook.com/home.php`
 - Custom HTTP header
 - Set custom headers for each REST request
 - Attacker cannot set custom header by script via form, image, iframe, etc.
 - Unless using JavaScript XMLHttpRequest or Flash
 - JavaScript same-origin policy prevent cross-site requests
 - Verify request's header contains X-Requested-By: XMLHttpRequest or X-Requested-With...
 - If no header, drop request
 - Recommendations
 - Strict referer validation for login forms or bank sites (info submitted over HTTPS)
 - For other sites, use Ruby-on-Rails or other framework that implements secret validation token (correctly)
 - Another type of header?
- Cross-Site Scripting (CSS/XSS)
 - Attacker injects malicious code into link to (supposedly) trustworthy source (sent to Alice via email, etc)
 - When user visits vulnerable site, embedded script is submitted as part of client's Web request (i.e. Google search)
 - If server-side application reflects user input (CSS vulnerability exists), browser will run reflected script
 - Bypasses same-origin policy test
 - Attacker can retrieve user authentication cookie and learn sensitive data, or hack web application itself
 - Defenses
 - HTML-escape all user input
 - Browser displays but does not run HTML-escaped input
 - Sanitize input by stripping of tags
- SQL Injection
 - Input validation vulnerability
 - User input in HTTP GET request could contain termination of line followed by (malicious) SQL script
 - E.g. `' ; DROP TABLE USERS; --`
 - Eliminates all user accounts
 - Prevention
 - Input validation
 - Filter characters with special meanings
 - Check data types
 - Whitelisting
 - Blacklisting "bad" characters doesn't work

- Could forget to filter out some characters
 - Could reject valid input
 - Allow only well-defined set of safe values
 - Implicitly defined through regular expressions
- Limit privileges
 - Prevent leakage of database schema
 - Encrypt sensitive data stored in database

Web Privacy

- The market for software that respect user privacy is a *lemons market*
- “Third party” online tracking
 - Sites other than the one you are visiting tracking your browsing history
 - Typically invisible to users
 - 64 independent tracking mechanisms on typical top-50 sites
- Tracking techniques
 - Tagging
 - Placing data in your browser
 - Includes: HTTP Cookies, HTTP Auth, HTTP Etags, Content cache, IE userdata, HTML 5 protocol & content handlers, HTML5 Storage, Flash cookies, Silverlight storage, TLS session ID & resume, Browsing history, window.name, HTTP STS, DNS cache
 - Fingerprinting
 - Observing your browser’s behavior
 - Includes: User-Agent, HTTP ACCEPT headers, Browser plugins, MIME support, Clock skew, installed fonts, cookies enabled?, browser add-ons, screen resolution
 - Browsers are unique enough
 - User agent string, plugins, etc. can uniquely identify
 - Panopticlick
 - Browser fingerprinting service/experiment
 - User-agent string: 10 bits of entropy, 84% of fingerprints unique (with Flash/Java, 94% unique)
- Anonymity?
 - Not quite
 - Third party is sometimes a first party
 - Facebook may be a third party to the site you are visiting, but if its “like” button is on the page, Facebook knows...(?)
 - Leakage of identifiers
 - GET `http://ad.doubleclick.net/adj/...`
Referer: `http://submit.SPORTS.com/...?email=jdoe@email.com`
Cookie: `id=35c192bcfe0000b1...`
 - If the email appears in the referer, identity has been compromised now and in the future
 - Third party buys your identity

- Hacks and bugs
 - Google spreadsheet (see github notes)
 - Cookie synchronization
 - Third party X sends its cookie to third party Y
 - X and Y exchange data about user
 - GET `http://tracker2.com/?uid=ghaihtn3`
Referer: `http://tracker1.com/...`
Cookie: `id=35c192bcfe0000b1...`
- Pseudonymity
 - Can tell when same person comes back (to website, etc) but doesn't know real-life identity
 - This is not true anonymity
 - Anonymity: shouldn't be able to track you under a pseudonym in a different session
 - Possible to connect online pseudonym with real-life identity
- Solutions
 - Referer blocking
 - Drawback: many sites check referer header for CSRF defense
 - Blocking referer indiscriminately will break sites
 - Drawback: can't prevent cooperative tracking
 - Third party cookie blocking
 - Advantage: does not break security systems
 - Drawback: doesn't prevent fingerprinting
 - Safari blocks third party cookies unless:
 - User is submitting a form
 - Browser already has cookie from same party
 - Do Not Track
 - Preference that can be set in web browsers
 - HTTP Request blocking
 - Compile and maintain list of known trackers
 - Semi-automated analysis
 - Based on domains and regular expressions
 - Sequence of events
 - User installs browser extension
 - Downloads list
 - Block request to objects on the list
 - Drawback: false positives and false negatives
 - Drawback: need to trust list
 - Blocking tools
 - Ghostery
 - Adblock Plus
 - Drawback: doesn't work by default – user must install
 - Drawback (Adblock): user needs to install blocklist separately to block all trackers like analytics and social widgets

Electronic Voting

- Types of voting machines
 - Hand-counter paper, punch cards, lever machines, optical scan ballots, electronic voting machines, touch-screen terminals, hybrid schemes
- Paper ballot attacks
 - Chain voting
 - Attacker obtains blank ballot and stands at entrance of voting booth
 - Attacker marks ballot as desired
 - Intimidates voter to take marked ballot and deposit it
 - Instructs voter to bring back blank ballot to attacker outside
 - Allows attacker to continue process
- “Receipt-free” secret ballot
 - Key aspect: cannot prove to 3rd party how you voted
- Proxy re-encryption
 - Bob wants to reveal contents of message sent to him (encrypted with his public key) to Chris without revealing private key to Chris
 - Designates proxy to re-encrypt message
 - Generates new key that Chris can use to decrypt message
 - Proxy cannot read Bob’s messages
- ElGamal encryption
 - Asymmetric key encryption algorithm based on Diffie-Hellman key exchange
 - Typically used in hybrid cryptosystem
 - Message encrypted using symmetric cryptosystem
 - ElGamal used to encrypt *key* used for symmetric cryptosystem
 - Algorithm
 - Each user has a private key x
 - Each user has three public keys: prime modulus p , generator g , and public $Y = g^x \bmod p$
 - Performance
 - As an asymmetric scheme, ElGamal is quite slow
 - Probabilistic
 - Advantage: single plaintext can be encrypted to many possible ciphertexts
 - Disadvantage: produces 2:1 expansion in size from plaintext to ciphertext
 - Security
 - Rests on the difficulty of the discrete log problem
 - Unconditionally malleable (not resistant to chosen ciphertext attack)
- End-to-end verifiability
 - Voter can confirm that vote was 1) cast as intended and 2) counted as cast
 - Does not have to trust election equipment or personnel
 - Should still be a secret-ballot
 - Goal: end-to-end verifiable elections while protecting voter privacy
 - Plan: use reencryption mix scheme
 - Two phases

- Voters publish their names and encrypted votes
 - Public “bulletin board” of ciphertext ballots
 - At end of election, administrators publish tally of votes
 - Include cryptographic proof that tally matches set of (published) encrypted votes
- Two possible paradigms
 - Anonymized ballots (mix networks)
 - Ballotless tallying (homomorphic encryption)
 - Includes RSA, ElGamal, Benaloh, etc.

Email Protocols

- Traditional mechanism
 - User composes message using email client on computer
 - Headers: to, from, date
 - Body: can encode different types of media
 - User hits “send” button
 - Email text and attachments uploaded to Simple Mail Transfer Protocol (SMTP) server as outgoing mail
 - Outgoing messages wait in outgoing mail queue
 - SMTP server communicates with DNS to find location of recipient’s email server
 - Messages are downloaded from recipient server to recipient’s email client
 - Uses Internet Message Access Protocol (IMAP)
 - Examples: Thunderbird, Postbox, Outlook (desktop clients)
- Webmail
 - Uses HTTP(S) to upload messages to sender’s mail server and to download from recipient’s mail service
 - Still use SMTP to transfer mail from sender to receiver servers
 - Examples: Gmail, Yahoo! Mail, AOL Mail

Anonymous Communication

- Is Internet anonymous?
 - No. IP addresses necessary for routing
 - Best case: pseudonymous
 - Worst case: identified
 - Encryption does not hide identities
- Tor (Onion Router)
 - Internet networking protocol designed to anonymize data relayed across it
 - Protects against Internet surveillance form known as “traffic analysis”
 - Data bundled into encrypted packet when it enters Tor network
 - Unlike normal internet connections, Tor
 - Strips away part of packet header
 - Separates addressing information that could identify sender
 - Encrypts rest of addressing information (packet wrapper)
 - Modified/encrypted data packet routed through many relays

- Each relays decrypts only enough of each data packet wrapper to know which relay data came from/which relay to send it to next
- Encryption keys different for each hop along circuit
- Last hop, from exit node to receiver, usually not encrypted
 - Cannot assume receiver is using Tor
- Goal: should not be able to trace data packet's path through Tor

SSH

- Cryptographic network protocol for secure data communication
 - Typically used to log into a remote machine and execute commands
- Connects, via secure channel over insecure network, a server (running SSH server) and a client (running SSH client)
- Uses public-key crypto to authenticate remote user
 - SSH only verifies whether public/private key match
 - Does not match public keys to identities
 - For unknown public keys, must verify this
- List of authorized public keys stored in Unix home directory
 - File located at ~/.ssh/authorized_keys
 - File should not be writable by anything apart from owner and root
 - SSH remembers key used by a server side over different sessions

Malware

- Taxonomy

	Requires host	Runs independently
Doesn't spread	Trojan, Rootkit	Keylogger, Spyware
Spreads	Virus	Worm

- Viruses
 - Definition: Reproduces own code by attaching itself to other executable files
 - When infected executable file is executed, virus code is also executed
 - Key points: self-replicating (spreads), infects files (requires host)
 - Classic viruses account for only 3% of all malware
 - What can act as a host?
 - Executable files
 - Either append code to file or overwrite parts of file code
 - Often take same name as existing files, with .exe extension
 - User might accidentally click and execute virus code
 - Boot sector
 - Region of hard disk containing machine code to be loaded into RAM on computer boot
 - E.g. Pakistani Brain virus
 - Macros
 - Set of instructions within application used to automate tasks

- Macros can perform system operations, such as creating, writing to, deleting files (potential for great damage)
 - Most macros written for Word, Excel, etc
 - Macro viruses infect templates for new documents
 - Each time new document is created, virus replicates
 - Cross-platform (not PC only)
- Virus lifecycle
 - Reproduction phase
 - Balances infection rate versus detection possibility
 - Infection phase
 - Viruses can stay resident in memory (dormant)
 - Attack phase
 - Attack on trigger
 - Jerusalem virus attacked on Friday the 13th
 - Delete files, change random data on disk
- Defenses
 - Antivirus software
 - Signature-based detection
 - Database of byte-level or instruction-level signatures that match virus (with wildcards, regular expressions)
 - Heuristics
 - Code execution starts in last section
 - Patched import address table
 - Sandboxing
 - Run untrusted applications in restricted environment
 - Default: do not run as administrator
- Variants
 - Encryption
 - Malware body encrypted with key
 - Decryption routine stored unencrypted
 - Decrypts upon execution
 - Polymorphic viruses
 - Change slightly with each infection
 - Encrypted payload
 - Different key used for each infection
 - Makes static string analysis difficult (impossible)
 - Metamorphic viruses
 - Different “versions” of code, but essentially same behavior
- Worms
 - Definition: self-replicating program that propagates itself across networks
 - Key points: self-replicating (spreads), propagates itself (no host)
 - Components
 - Target locator
 - Email harvesting (scan address books, inbox of email client, Google searches, buy list of emails, IP addresses)

- Infection propagator
 - Life cycle manager
 - Payload
 - Often a Trojan horse
- Variants
 - Email-based
 - Forged from address
 - Hide executable extension (.exe) behind harmless ones (.jpeg)
 - Promise interesting pictures or applications
 - Exploit-based
 - Do not require human interaction
 - Spread using well-known network services (TCP, etc)
 - Spread can be modeled with classic disease model
 - Slow start, followed by exponential growth
- Defenses
 - Virus scanners
 - Scan email attachments or other contents
 - Effective against email-based worms
 - Host level defense
 - Elimination of underlying software vulnerabilities
 - StackGuard: protect against buffer overflow
 - Randomize position of stack, heap, libraries in memory
 - Network level defense
 - Intrusion detection systems
 - Scan for known attack patterns
 - Rate limiting
 - Quota on number of outgoing connections
 - Personal firewall
 - Block outgoing SMTP connections from unknown apps
- Writing secure code
 - Careful coding, code audits, high-level languages, model checking, formal methods and protocol verification, fuzz testing, static analysis, dynamic analysis, taint analysis, comparison across implements, access control

Big Data and Privacy

- Goals
 - Make valid inferences about population as a whole from dataset
 - Cannot make valid inferences about individuals from dataset
- Semantic privacy
 - Given two datasets D and D' , where D' is D with one datapoint removed, anything analyst can learn from D , they can also learn from D'
 - Theorem: Semantic privacy implies result of analysis does not depend on content of dataset
- Differential privacy
 - Property of a protocol A run on a dataset X producing output $A(X)$

- A is a randomized algorithm
 - A gives ϵ -differential privacy if $A(X)$ and $A(X')$ give very similar results, where X and X' differ in the inclusion/exclusion of one element
- Post-processing
 - Theorem: applying an arbitrary function f to the output of a differentially private protocol A gives an output that is still ϵ -DP
- Achieving differential privacy
 - Output perturbation: add random noise to true answer of query
 - Use Gaussian distribution or (better) Laplace/geometric distribution
- Applications
 - Collaborative recommendation systems
 - “People who bought X also bought Y ”
 - Privacy issues
 - Rare book X only Ed would buy
 - Collaborative recommendation links item X to another item Y
 - Recommendation gives hint about what (else) Ed bought
- Solutions
 - Look at algorithm internals
 - System generates covariance matrix
 - Correlation between purchases of all pairs of items
 - Add random noise to matrix to achieve differential privacy
 - Machine learning and DP queries
 - Machine learning algorithm exchanges DP queries and results
 - Can synthesize new dataset

Economics and Security

- Fundamental question
 - Does the market produce optimal security?
- Definitions of “optimal”
 - Strong Pareto Efficiency
 - Condition A is SP-superior to Condition B if everyone prefers A over B
 - Condition is SP-efficient if no SP-superior alternative available
 - Impossible to make any one individual better off with making at least one individual worse off
 - Pareto improvement: a change that could make one individual better off without making any other individual worse off
 - Kaldor-Hicks Efficiency
 - Condition A is KH-superior to Condition B if a set of zero-sum payments P among people (i.e. wealth transfers, redistributive taxes) exists such that $A + P$ is SP-superior to Condition B
 - Note: payments need not occur in practice
 - Condition is SP-efficient if no KH-superior alternative is available
 - Theorem: a world with perfect information and perfect bargaining would be SP-efficient and KH-efficient
 - Proof by contradiction

- Implication: since world isn't SP-efficient or KH-efficient, market failures must be occurring
- Market failure #1: negative externalities
 - Harm falls on third party (not seller or buyer)
 - Neither will invest in reducing harm to third party
 - Implication: underinvestment in security
 - Note: bargaining to fix externalities not possible in real world
- Market failure #2: asymmetric information
 - Hard for buyers to evaluate security of products
 - Producer knows more about security of product than buyer
 - "Lemons market"
 - Little incentive for producer to improve quality
 - Solutions
 - Add warranties to product
 - Seller reputation
- Network effects
 - Product becomes more valuable as more people use it
 - Tends to push markets toward monopoly (monoculture)
 - Benefits of having a dominant producer
 - Security is often more efficient with scale
 - As a producer whose product pervades society, some of the external benefits are in fact internalized (no longer true externalities)
 - Warranties and reputation matter more
 - Nuance: race to market
 - Network effects often tip toward early leader
 - Companies try to get MVP (minimum viable product) into market as soon as possible
 - Less incentive to work on security now
 - "Bolt on security later" approach
 - Solutions
 - Large customers can protect themselves
 - Market structures to improve information flow
 - Insurance companies, certification programs
 - Change in liability rules
 - Optimal rule: cost born by whoever can best prevent harm (this tends to be the producer)
 - Problems: hard to attribute blame, hard to measure harm, and high cost to adjudication (judging)
 - Public inspections
 - Large buyer demands ability to publicize security evaluations of products

Human Factors in Security

- Reasons for user error
 - Bad UI/UX leads to mistakes

- If pilot makes mistake, system should change to make that mistake harder to make (blame system, not person in long run)
 - Rational ignorance
 - Reason: security/system is too difficult to understand
 - Cost of user informing him/herself seems higher than cost of breach
 - Heuristic decision making and cognitive biases
 - Could be exploitable by adversary
 - Relying on user intelligence/designing for yourself
- Wifi encryption
 - General recommendation: wifi networks should be encrypted
 - Reality: open wifi networks are not encrypted
 - PUK Wireless is a closed network that should be encrypted, but isn't
 - Problem
 - Key distribution to all devices using wifi network
 - Someone joins airport wifi access point to access internet, but doesn't know how to enter key
 - Possible solutions
 - Exploit physical proximity between devices
 - "Tap to pair this device"
 - Line-of-sight medium
 - Trust on first use policy (TOFU)